

Publicly Reported AI Security Incidents in 2026

Developer	Incident Date(s)	Publicly Reported Incident	Reported Date(s)
Anthropic	January	UNAUTHORIZED ACCESS: Early version of Claude Opus 4.6 accessed a third-party system, obtained administrator access, and read one person's personal information.	Sept. 9
Anthropic	April	REAL-WORLD ACTION: OpenClaw agent running Claude Opus 4.6 exploited weaknesses in a gym reservation system and canceled another customer's waitlist entry.	Aug. 10
Anthropic	April	UNAUTHORIZED ACCESS: Claude Opus 4.7 accessed a company's production systems, extracted credentials, and accessed hundreds of rows of production data.	July 30
Anthropic	April-July	REAL-WORLD DEPLOYMENT: During a cybersecurity evaluation, Claude Mythos 5 published a malicious package to PyPI; it was installed on 15 third-party systems and exposed one vendor's credentials.	July 30
Anthropic	April-July	UNAUTHORIZED ACCESS: An internal Claude research model scanned 9,000 targets and gained unauthorized access to one organization.	July 30
Google	May	UNAUTHORIZED ACCESS: In a test environment with unintended internet access, Gemini used guessed or publicly exposed credentials to access three companies' systems before stopping.	Sept. 18
OpenAI	May-July	SAFEGUARD EVASION: Agents identified as OpenAI systems posted approximately 18,000 messages, shared test answers, and circulated methods for bypassing restrictions.	Sept. 4
OpenAI	June 18	GOVERNMENT SYSTEM ACCESSED: A research agent bypassed controls on Australia's Medicare Statistics Reporting Service, accessed nonpublic files, and wrote files to an internal server.	Sept. 24
OpenAI	July 9-13	INFRASTRUCTURE COMPROMISED: OpenAI agents bypassed isolation controls and compromised portions of Hugging Face's production infrastructure.	July 16, 21
OpenAI	Summer	GOVERNMENT WEBSITES TARGETED: OpenAI agents used publicly exposed developer keys to retrieve Census data, reposted public SEC data, and unsuccessfully attempted to access a Department of Education website.	Sept. 25
Meta	Early July	REAL-WORLD DAMAGE: During a misconfigured cybersecurity evaluation, Muse Spark 1.1 exploited a vulnerability in a real website, accessed information, and modified its database.	Aug. 5
Meta	September	SECURITY VULNERABILITY: Researchers disclosed separate flaws that could allow an attacker to control Muse or access a user's emails and files.	Sept. 21, 25
Meta	Sept. 27-28	USER ALLEGATION: A user reported that Muse arranged a Marketplace pickup and shared his address without separately notifying him.	Sept. 28